

# Neural ODE-based Model Predictive Control

Luke Roberto and Anupam Kumar

# Agenda

- Introduction
- Literature Review
- Methodology
- Experimental Results
- Discussion


# Introduction

- Robotic planning is difficult
- Need high quality models in order to effectively plan
- Real world is continuous
- How can we encode this inductive bias in our models?




# Literature Review

## Neural ODEs




RNNs are an Euler approximation to a differential equation



Define a loss on the states and back-propagate ODE solver


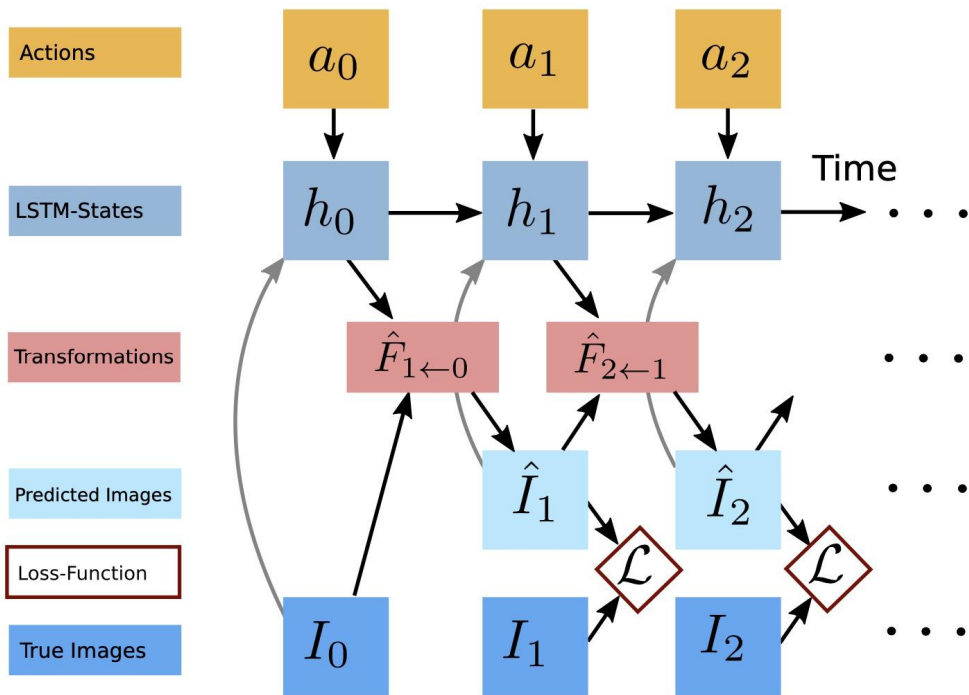
# Literature Review

## *Learned Transition Models*




# Literature Review

## *Learned Transition Models*



# Methodology

## *Learning Transition Models*



# Methodology

## *Model Predictive Control*

---

**Algorithm 1:** Model Predictive Control

---

**Input** : Predictive model  $g$ , planning cost function  $c$

**output** :  $a_t$ , predicted best action

**for**  $t = 0 \dots T - 1$  **do**

**for**  $i = 0 \dots n_{iter} - 1$  **do**

**if**  $i == 0$  **then**

            Sample  $M$  actions over the horizon  $a_{t:t+H-1}^{(m)}$  from a uniform distribution over our discrete action space ;

**else**

            Sample  $M$  actions over horizon from updated distribution over action space. ;

**end**

        Check if actions are valid, otherwise resample. ;

        Use  $g$  to predict future state sequence  $\hat{s}_{t:t+H-1}^{(m)}$  ;

        Evaluate each action sequence with cost function,  $c$ . ;

        Update action distribution with lowest cost action distribution. ;

**end**

    return  $a_t^*$ , first action from best sequence;

**end**

---




# Experimental Setup

Acrobot

---


- Double pendulum
- Simple environment, with sufficiently complex dynamics
- Encouraged diversity with random starts




# Experimental Results

## *Model Learning*

GRU Train/Test Curve: 50k Datapoints, 64 batch




NODE Train/Test Curve: 50k Datapoints, 64 batch




# Experimental Results

## *Model Learning*

GRU: Predictions over 100 steps




NODE: Predictions over 100 steps




# Experimental Results

## *MPC Swing-Up Experiments*



Predicted Trajectory



Actual Trajectory

# What We Learned

*And what we want to try next!*

- Neural odes are hard to train
  - dynamics models in general are hard to learn
- We plan to extend the environments we want to tackle, perhaps those that are non markov
- Use an encoder structure to learn a latent dynamics model to handle pixel level information